

Computer Organization and Design

Instruction Sets - 1

Lecture 4

Representing Instructions

* Today's topics

- von Neumann model of a computer
- Instruction set architecture
- MIPS instruction formats
- Some MIPS instructions

* Reading

- P&H textbook Ch. 2.1-2.2, Ch. 2.5-2.6

Levels of Program Code

* High-level language

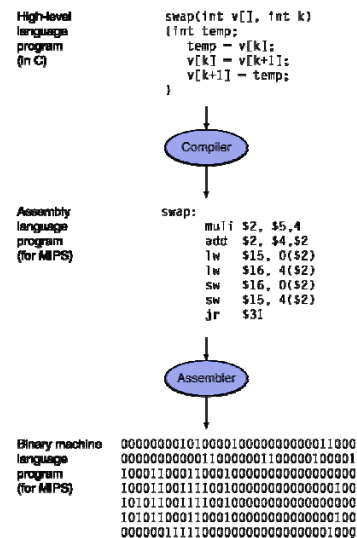
- Level of abstraction closer to problem domain
- Provides for productivity and portability

* Assembly language

- Textual representation of instructions

* Hardware representation

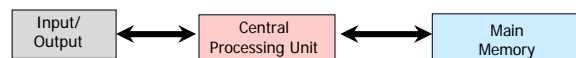
- Binary digits (bits)
- Encoded instructions and data



A General-Purpose Computer

* The von Neumann Model

- Many architectural models for a general-purpose computer have been explored
- Most of today's computers based on the model proposed by John von Neumann in the late 1940s
- Its major components are:

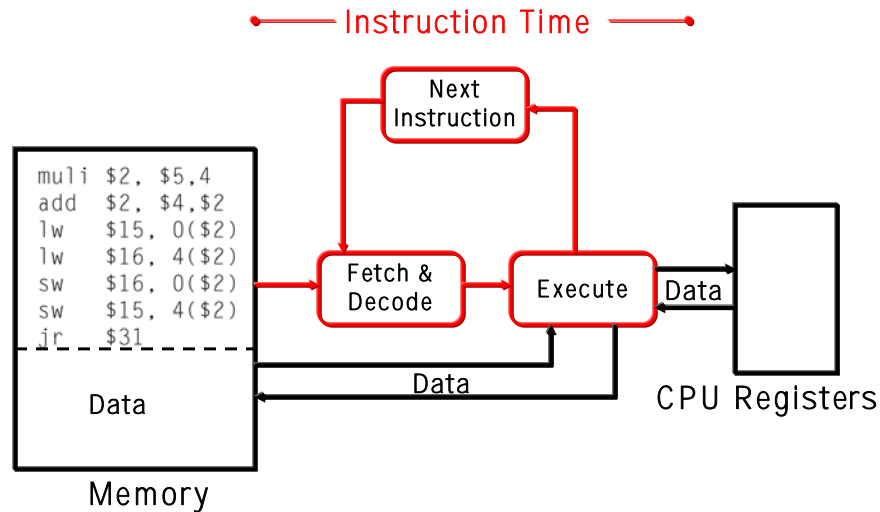


Central Processing Unit (CPU): Fetches, interprets, and executes a specified set of operations called **Instructions**.

Memory: storage of N *words* of W bits each, where W is a fixed architectural parameter, and N can be expanded to meet needs.

I/O: Devices for communicating with the outside world.

CPU (based on Von Neumann)



5

Instructions and Programs

* What are *instructions*?

- the words of a computer's language

* Instruction Set

- the full vocabulary

* Stored Program Concept

- The idea that instructions and data of many types can be stored in memory as numbers, leading to the stored program computer
 - Distinct from "application-specific" hardware, which is "hardwired" to perform "fixed-function" processing on inputs
 - Distinct from punched tape computers (e.g., looms) where instructions were not stored, but streamed in one at a time

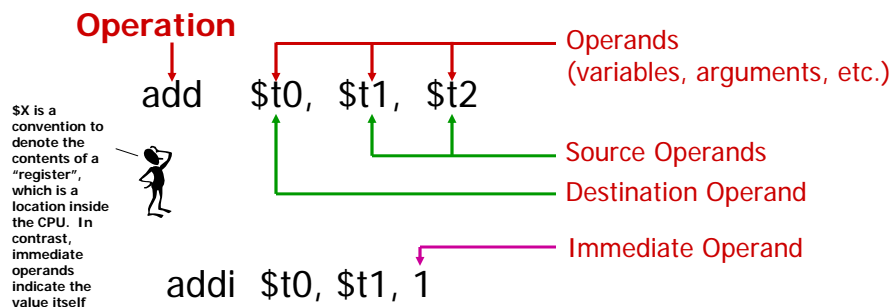
Our representative example: the **MIPS** instruction set

6

Anatomy of an Instruction (Assembler)

* An instruction is a primitive operation

- Instructions specify an operation and its operands (the necessary variables to perform the operation)
- Types of operands: immediate, source, and destination



7

Meaning of an Assembler Instruction

* Operations are abbreviated into opcodes (1-4 letters)

* Instructions are specified with a very regular syntax

- First an opcode followed by arguments
- Usually (but not always) the destination is next, then source
- Why this order? Arbitrary...
 - ... but analogous to high-level language like Java or C

`add $t0, $t1, $t2`

↓ implies

`$t0 = $t1 + $t2`

The instruction syntax provides operands in the same order as you would expect in a statement from a high level language.

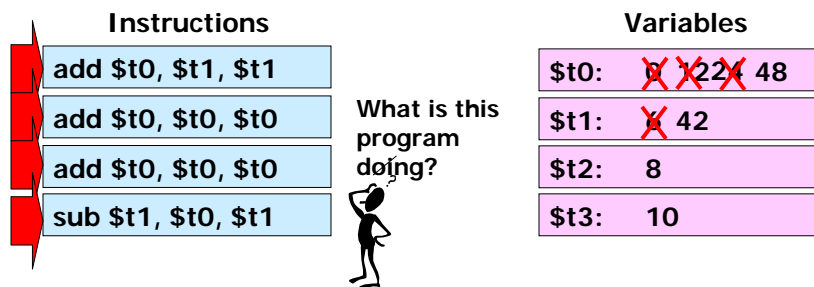


8

Being the Machine!

* Instruction sequence


- Instructions are executed sequentially from a list ...
 - ... unless some special instructions alter this flow
- Instructions execute one after another
 - therefore, results of all previous instructions have been computed



9

What did this machine do?

- * Let's repeat the simulation, this time using unknowns
- * Knowing what the program does allows us to write down its specification, and give it a meaningful name

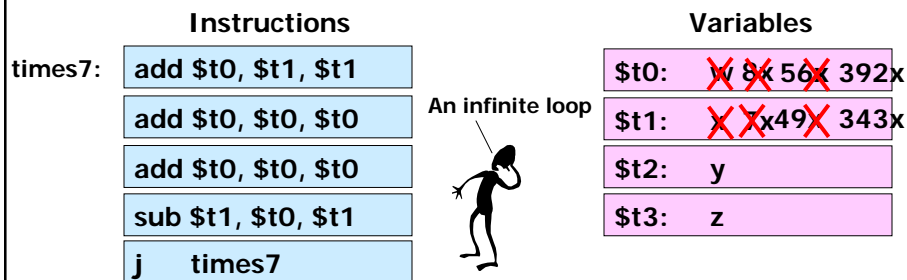


10

Looping the Instruction Sequence

* Need something to change the instruction flow

- “go back” to the beginning
- a *jump* instruction with opcode ‘j’
 - the operand refers to a label of some other instruction
 - for now, this is a text label you assign to an instruction
 - in reality, the text label becomes a numerical address



11

Open Questions in our Simple Model

* We will answer the following questions next

- WHERE are INSTRUCTIONS stored?
- HOW are instructions represented?
- WHERE are VARIABLES stored?
- How are labels associated with particular instructions?
- How do you access more complicated variable types:
 - Arrays?
 - Structures?
 - Objects?
- Where does a program start executing?
- How does it stop?


12

The Stored-Program Computer

* The von Neumann model:

- Instructions and Data stored in a common memory ("main memory")
- Sequential semantics: All instructions execute sequentially (or at least appear sequential to the programmer)

Key idea: Memory holds not only data, but *coded instructions* that make up a *program*.

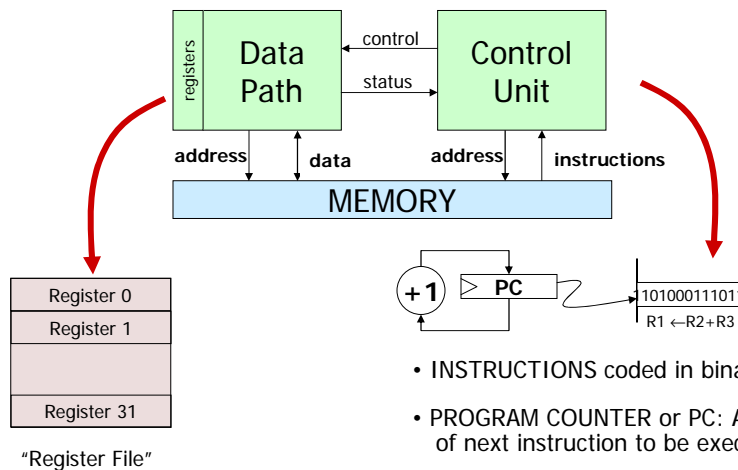


* CPU fetches and executes instructions from memory ...

- The CPU is a H/W interpreter
- Program **IS** simply data for this interpreter
- Main memory: Single expandable resource pool
 - constrains both data and program size
 - don't need to make separate decisions of how large of a program or data memory to buy

13

More details of a von Neumann CPU




- INSTRUCTIONS coded in binary
- PROGRAM COUNTER or PC: Address of next instruction to be executed
- Control Unit has circuitry inside to translate instructions into control signals for data path

14

The big picture

* A few things to note:

- Memory is distinct from data path
- Registers are in data path
- Program is stored in memory
- Control unit fetches instructions from memory
- Control unit tells data path what to do
- Data can be moved from memory to registers, or from registers to memory
- All data processing (e.g., arithmetic) takes place within the data path



15

Instruction Set Architecture

* Definition:

- The part of the computer architecture related to programming, including the native data types, instructions, registers, addressing modes, memory architecture, interrupt and exception handling, and external I/O
- An ISA includes a specification of the set of opcodes (machine language), and the native commands implemented by a particular processor

16

Instruction Set Architecture (ISA)

* Encoding of instructions raises interesting choices...

- Tradeoffs: performance, compactness, programmability
- Complexity
 - How many different instructions? What level operations?
 - Level of support for particular software operations: array indexing, procedure calls, “polynomial evaluate”, etc.
 - “Reduced Instruction Set Computer” (RISC) philosophy: simple instructions, optimized for speed
- Uniformity
 - Should different instructions be same size?
 - Take the same amount of time to execute?
 - Trend favors uniformity → simplicity, speed, cost/power

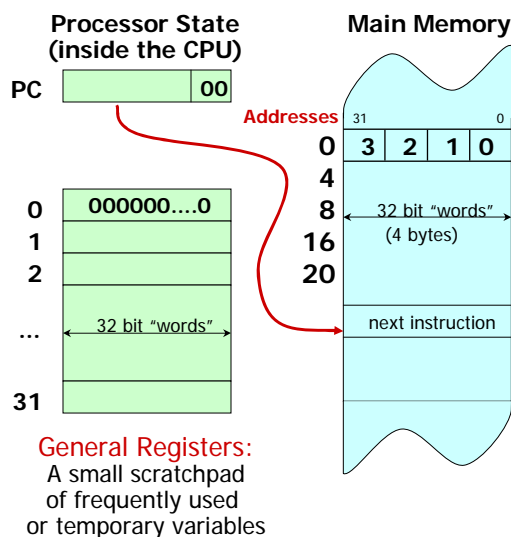
* Mix of Engineering & Art...

- Trial (by simulation) is our best technique for making choices!

17

MIPS Programming Model

a representative simple RISC machine



In Comp 411 we'll use a clean and sufficient subset of the MIPS-32 core Instruction set.

Fetch/Execute loop:

- fetch Mem[PC]
- $PC = PC + 4^\dagger$
- execute fetched instruction (may change PC!)
- repeat!

[†]MIPS uses byte memory addresses. However, each instruction is 32-bits wide, and *must* be aligned on a multiple of 4 (word) address. Each word contains four 8-bit bytes. Addresses of consecutive instructions (words) differ by 4.

18

Some MIPS Memory Details


* Memory locations are 32 bits wide

- Addressable in different-sized chunks specified by instruction
- 8-bit chunks (byte)
- 16-bit chunks (halfword)
- 32-bit chunks (word)
- 64-bit chunks (doubleword)

* We also frequently need access to individual bits! (Instructions help w/ this)

* Every BYTE has a unique address (MIPS is a byte-addressable machine)

* Every instruction is one word



19

MIPS Register Details

* There are 32 named registers [\$0, \$1, \$31]

* The operands of all ALU instructions are registers

- This means to operate on a variables in memory you must:
 - Load the value/values from memory into a register
 - Perform the instruction
 - Store the result back into memory
- Going to and from memory can be expensive
 - (4x to 20x slower than operating on a register)
- Net effect: Keep variables in registers as much as possible!

* Special purpose and conventions

- 2 registers have specific "side-effects"
 - (ex: \$0 always contains the value '0'... more later)
- 4 registers dedicated to specific tasks by convention
- 26 available for general use, but constrained by convention

20

MIPS Register Names and Conventions

- * Some MIPS registers assigned to specific uses

- E.g., \$0 is hard-wired to the value 0; \$31 used for subroutine linkage

- * Each register has two symbolic names

- E.g., \$zero and \$0, \$t0 and \$8 (see table below)

Name	Register number	Usage
\$zero	0	the constant value 0
\$at	1	assembler temporary
\$v0-\$v1	2-3	values for results and expression evaluation
\$a0-\$a3	4-7	arguments
\$t0-\$t7	8-15	temporaries
\$s0-\$s7	16-23	saved
\$t8-\$t9	24-25	more temporaries
\$gp	28	global pointer
\$sp	29	stack pointer
\$fp	30	frame pointer
\$ra	31	return address

21

MIPS Instruction Formats

- * All MIPS instructions fit into a single 32-bit word

- * Every instruction includes various "fields":

- a 6-bit operation or "OPCODE"
 - specifies which operation to execute (fewer than 64)
- up to three 5-bit OPERAND fields
 - each specifies a register (one of 32) as source/destination
- embedded constants
 - also called "literals" or "immediates"
 - 16-bits, 5-bits or 26-bits long
 - sometimes treated as signed values, sometimes unsigned

- * There are three basic instruction formats:

- **R-type**, 3 register operands (2 sources, destination)
- **I-type**, 2 register operands, 16-bit constant
- **J-type**, no register operands, 26-bit constant

OP	r_s	r_t	r_d	shamt	func
----	-------	-------	-------	-------	------


OP	r_s	r_t	16-bit constant
----	-------	-------	-----------------

OP	26-bit constant
----	-----------------

22

MIPS ALU Operations

Sample coded operation: ADD instruction



References to register contents are prefixed by a '\$' to distinguish them from constants or memory addresses



What *we* prefer to write: `add $10, $11, $9` ('assembly language')

The convention with MIPS assembly language is to specify the destination operand first, followed by source operands.



add rd, rs, rt:

$\text{Reg}[\text{rd}] = \text{Reg}[\text{rs}] + \text{Reg}[\text{rt}]$

"Add the contents of rs to the contents of rt; store the result in rd"

Similar instructions for other ALU operations:

arithmetic: add, sub, addu, subu


compare: slt, sltu

logical: and, or, xor, nor

shift: sll, srl, sra, sllv, srav, srlv

23

Programming without Compilers/Assemblers



Shift operations

* Shifting is a common operation


- applied to groups of bits
- used for alignment
- used for "short cut" arithmetic operations
 - $X \ll 1$ is often the same as $2 * X$
 - $X \gg 1$ can be the same as $X / 2$

* For example:

- $X = 20_{10} = 00010100_2$
- Left Shift (Logical):
 - $(X \ll 1) = 00101000_2 = 40_{10}$
- Right Shift (Logical):
 - $(X \gg 1) = 00001010_2 = 10_{10}$
- Signed or "Arithmetic" Right Shift:
 - $(-X \gg 1) = (11101100_2 \gg 1) = 11110110_2 = -10_{10}$

MIPS Shift Operations

Sample coded operation: SHIFT LOGICAL LEFT instruction



Assembly: `sll $2, $2, 4`


`sll rd, rt, shamt:`

`Reg[rd] = Reg[rt] << shamt`

*Shift the contents of *rt* to the left by *shamt*; store the result in *rd**

MIPS Shift Operations

Sample coded operation: SLLV (SLL Variable)



This is peculiar syntax for MIPS. In this ALU instruction the rt operand precedes the rs operand. Usually, it's the other way around



Different flavor:
Shift amount is not in instruction, but in a register

Assembly: `sllv $2, $2, $8`

`sllv rd, rt, rs:`


$\text{Reg}[rd] = \text{Reg}[rt] \ll \text{Reg}[rs]$

*Shift the contents of *rt* left by the contents of *rs*; store the result in *rd**

27

MIPS ALU Operations with Immediate

addi instruction: adds register contents, signed-constant:



Symbolic version: `addi $9, $11, -3`

`addi rt, rs, imm:`

$\text{Reg}[rt] = \text{Reg}[rs] + \text{sxt}(imm)$

*Add the contents of *rs* to const; store result in *rt**

Similar instructions for other ALU operations:

arithmetic: `addi, addiu`
compare: `slti, sltiu`
logical: `andi, ori, xori, lui`

Immediate values are sign-extended for arithmetic and compare operations, but zero extended for logical operations.



28

Why Built-in Constants? (Immediate)

* Where are constants/immediates useful?

- SMALL constants used frequently (50% of operands)

- In a C compiler (gcc) 52% of ALU operations use a constant
- In a circuit simulator (spice) 69% involve constants
- e.g., $B = B + 1$; $C = W \& 0x00ff$; $A = B + 0$;

* Examples:

```
addi $29, $29, 4
slti $8, $18, 10
andi $29, $29, 6
ori  $29, $29, 4
```

29

First MIPS Program (fragment)

* Suppose you want to compute the expression:

$$f = (g + h) - (i + j)$$

- where variables f, g, h, i, and j are assigned to registers \$16, \$17, \$18, \$19, and \$20 respectively
- what is the MIPS assembly code?

```
add $8,$17,$18      # (g + h)
add $9,$19,$20      # (i + j)
sub $16,$8,$9       # f = (g + h) - (i + j)
```

- Questions to answer:

- How did these variables come to reside in registers?
- Answer: We need more instructions which allow data to be explicitly loaded from memory to registers, and stored from registers to memory

30